

# Database Management System (DBMS)

## Lecture-44

---


Dharmendra Kumar

November 24, 2022

## Multiple Granularity

---

Consider the following granularity hierarchy.



# Concurrency Control

This tree consists of four levels of nodes. The highest level represents the entire database. Below it are nodes of type area; the database consists of exactly these areas. Each area in turn has nodes of type file as its children. Each area contains exactly those files that are its child nodes. No file is in more than one area. Finally, each file has nodes of type record. As before, the file consists of exactly those records that are its child nodes, and no record can be present in more than one file.

# Concurrency Control

This protocol uses the following compatibility matrix to lock the data items.

**Lock Compatibility Matrix**

	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false

There is an intention mode associated with shared mode, and there is one with exclusive mode. If a node is locked in intention-shared (IS) mode, explicit locking is being done at a lower level of the tree, but with only shared-mode locks.

Similarly, if a node is locked in intention-exclusive (IX) mode, then explicit locking is being done at a lower level, with exclusive-mode or shared-mode locks.

Finally, if a node is locked in shared and intention-exclusive (SIX) mode, the sub-tree rooted by that node is locked explicitly in shared mode, and that explicit locking is being done at a lower level with exclusive-mode locks.

# Concurrency Control

**The multiple-granularity locking protocol**, which ensures serializability, is this:

Each transaction  $T_i$  can lock a node  $Q$  by following these rules:

1. It must observe the lock-compatibility function shown in above matrix.
2. It must lock the root of the tree first, and can lock it in any mode.
3. It can lock a node  $Q$  in S or IS mode only if it currently has the parent of  $Q$  locked in either IX or IS mode.
4. It can lock a node  $Q$  in X, SIX, or IX mode only if it currently has the parent of  $Q$  locked in either IX or SIX mode.
5. It can lock a node only if it has not previously unlocked any node (that is,  $T_i$  is two phase).
6. It can unlock a node  $Q$  only if it currently has none of the children of  $Q$  locked.

# Concurrency Control

Clearly, the multiple-granularity protocol requires that locks be acquired in top-down (root-to-leaf) order, whereas locks must be released in bottom-up (leaf-to-root) order.

## Example:

Consider the tree shown in the above figure and these transactions:

- Suppose that transaction  $T_{18}$  reads record  $r_{a_2}$  in file  $F_a$ .  
Then,  $T_{18}$  needs to lock the database, area  $A_1$ , and  $F_a$  in IS mode (and in that order), and finally to lock  $r_{a_2}$  in S mode.
- Suppose that transaction  $T_{19}$  modifies record  $r_{a_9}$  in file  $F_a$ .  
Then,  $T_{19}$  needs to lock the database, area  $A_1$ , and file  $F_a$  in IX mode, and finally to lock  $r_{a_2}$  in X mode.

# Concurrency Control

- Suppose that transaction  $T_{20}$  reads all the records in file  $F_a$ . Then,  $T_{20}$  needs to lock the database and area  $A_1$  (in that order) in IS mode, and finally to lock  $F_a$  in S mode.
- Suppose that transaction  $T_{21}$  reads the entire database. It can do so after locking the database in S mode.

Clearly, transactions  $T_{18}$ ,  $T_{20}$ , and  $T_{21}$  can access the database concurrently. Transaction  $T_{19}$  can execute concurrently with  $T_{18}$ , but not with either  $T_{20}$  or  $T_{21}$ .



This protocol enhances concurrency and reduces lock overhead. It is particularly useful in applications that include a mix of

- Short transactions that access only a few data items
- Long transactions that produce reports from an entire file or set of files

**Note:** Deadlock is possible in this protocol.

## Multiversion Schemes

---

In multiversion concurrency control schemes, each  $\text{write}(Q)$  operation creates a new version of  $Q$ . When a transaction issues a  $\text{read}(Q)$  operation, the concurrency control manager selects one of the versions of  $Q$  to be read. The concurrency control scheme must ensure that the version to be read is selected in a manner that ensures serializability.

## Multiversion Timestamp Ordering

---

With each data item  $Q$ , a sequence of versions  $\langle Q_1, Q_2, \dots, Q_m \rangle$  is associated. Each version  $Q_k$  contains three data fields:

- **Content** is the value of version  $Q_k$ .
- **W-timestamp( $Q_k$ )** is the timestamp of the transaction that created version  $Q_k$ .
- **R-timestamp( $Q_k$ )** is the largest timestamp of any transaction that successfully read version  $Q_k$ .

A transaction  $T_i$  creates a new version  $Q_k$  of data item  $Q$  by issuing a write( $Q$ ) operation. The content field of the version holds the value written by  $T_i$ . The system initializes the W-timestamp and R-timestamp to  $TS(T_i)$ . It updates the R-timestamp value of  $Q_k$  whenever a transaction  $T_j$  reads the content of  $Q_k$ , and  $R\text{-timestamp}(Q_k) < TS(T_j)$ .

**The multiversion timestamp-ordering scheme** operates as follows. Suppose that transaction  $T_i$  issues a read(Q) or write(Q) operation. Let  $Q_k$  denote the version of Q whose write timestamp is the largest write timestamp less than or equal to  $TS(T_i)$ .

1. If transaction  $T_i$  issues a read(Q), then the value returned is the content of version  $Q_k$ .
2. If transaction  $T_i$  issues write(Q), and if  $TS(T_i) < R\text{-timestamp}(Q_k)$ , then the system rolls back transaction  $T_i$ . On the other hand, if  $TS(T_i) = W\text{-timestamp}(Q_k)$ , then the system overwrites the contents of  $Q_k$ ; otherwise it creates a new version of Q.

# Concurrency Control

Versions that are no longer needed are removed according to the following rule. Suppose that there are two versions,  $Q_k$  and  $Q_j$ , of a data item, and that both versions have a W-timestamp less than the timestamp of the oldest transaction in the system. Then, the older of the two versions  $Q_k$  and  $Q_j$  will not be used again, and can be deleted.

## Note:

1. The multiversion timestamp-ordering scheme ensures serializability.
2. The multiversion timestamp-ordering scheme does not ensure recoverability and cascadelessness.